

# A Case Study on Performance Optimization Techniques in Java Programming

Ciprian Khlud<sup>1</sup> <sup>a</sup> and Cristian Frăsinaru<sup>1</sup> <sup>b</sup>

<sup>1</sup>"Alexandru Ioan Cuza University", Iași, Romania  
ciprian.mustiata@gmail.com, acf@info.uaic.ro

**Keywords:** Java, Runtime performance, Memory usage, Garbage collection, Sequence analysis, SAM/BAM files

**Abstract:** Choosing the right programming platform for processor or memory intensive applications is a subject that is debated in all types of contexts. When analyzing the performance of a specific platform, equally important is the usage of appropriate language specific constructions and programming interfaces (APIs). In this paper we investigate how a state-of-the art implementation, part of a multi-threaded framework for sequence analysis (elPrep) could benefit from various optimization techniques dedicated to improving the runtime performance of Java applications. ElPrep is an established tool for processing SAM and BAM files in sequencing pipelines. In order to migrate from its original implementation to a different software platform, more suitable for memory intensive tasks, the authors have re-implemented elPrep in Java, Go and C++ and benchmarked their performance. Overall, the Go implementation won by a good margin, considering a metric that involved both the RAM usage and the runtime performance. We show that, without changing the semantics of the algorithm, by using appropriate programming techniques we are able to significantly improve the behavior of the Java implementation, to a point that may even alter the conclusions of the original study. We also show, that with changing the semantics of the algorithm, especially focusing on memory consumption, we are able to improve the original scoring (based of a composite of compute time and memory) to around one order of magnitude better on the most expensive component (read/write).


## 1 INTRODUCTION


In the field of bioinformatics, DNA sequence analysis generally consists of processing large amounts of data and performing various operations on it, such as sequence alignment, variant detection, searches against biological databases, etc. A large variety of software tools exists for these operations, most of them having specific uses cases but with a common denominator regarding the fact they need to perform processor and memory intensive tasks: I/O operations on large file, compression/decompression, text processing, etc. (?)

Choosing a programming platform that offers all the required instruments to handle the specific challenges in bioinformatics is important, as pointed out in a recent study dedicated to migrating an existing Common Lisp application, called elPrep, to another platform with better support for memory management and concurrency (Costanza et al., 2019). ElPrep (Herzeel et al., 2019) is a a multithreaded

tool for preparing sequence alignment/map files (SAM/BAM) for variant calling in DNA sequencing pipelines. A key feature of elPrep is the ability to avoid the standard practice of creating a pipeline consisting of multiple command line tools invocations, executing a single pass through a SAM/BAM file and keeping data as much as possible in main memory. In (Costanza et al., 2019) the authors investigated Go, Java and C++ programming platforms, as an alternative to Common Lisp. Te result of their study concluded that the Go implementation performed best, using a metric that involved both the RAM usage and the runtime performance. The benchmarks of the study showed that Java had a faster runtime, but a significantly higher memory usage, while Go offered a better balance between the two.

As the Java source code for elPrep is available at <https://github.com/exascience/elprep-bench>, we have analyzed key aspects regarding the memory management and thread synchronization, and propose a series of improvements that could increase significantly the performance of the Java

<sup>a</sup>  <https://orcid.org/0000-0001-6211-3199>

<sup>b</sup>  <https://orcid.org/0000-0002-5246-7396>

implementation.

## 2 BACKGROUND

### 2.1 GARBAGE COLLECTION

In order to analyze the behavior of memory intensive applications, it is important to understand how garbage collection works and especially how Java (Java Platform, Standard Edition, 2019) implements its garbage collectors.

The Java Virtual Machine (JVM) (Lindholm et al., 2014) offers an automatic storage management system, called *garbage collector* (GC) which reclaims heap storage occupied by objects which are no longer used. The garbage collection process (Oracle, 2019) works typically by splitting the heap into two regions: a *young generation* region and an *old generation*. All new objects are allocated in the young region, in a very fast manner, using typically a "bump-pointer" strategy. When this region becomes full a *minor* garbage collection occurs and all dead objects are deleted very quickly. The objects which are still referenced survive and they are moved to the old generation. This minor collection is always a "stop the world" event, meaning that all of the application threads will be paused until the GC is finished. In the old generation, objects are expected to live longer and they are collected more seldom but with a more expensive algorithm, called *major* garbage collection.

The algorithm used by GC has two steps. The first one is to *mark* the objects that are still used from the heap. In the second step, it *sweeps* (deletes) the objects which have not been marked (dead), leaving only referenced objects and pointers to free space. By moving referenced object together, this makes new memory allocation much easier and faster. Therefore, the speed of GC depends on two factors: the number of objects it has to analyze and the complexity of the relationships between them.

Considering the behavior we have described so far, we will analyze the impact of some simple tweaks meant to reduce the impact of GC over the application performance, such as: reducing the unnecessary small allocations in young region, controlling the scope in which objects are referenced in order to minimize the number of times when expensive collection of old region is triggered, simplifying the object graph and controlling the amount of memory JVM is allowed to use.

### 2.2 THREADS, LOCKS AND THE FILE SYSTEM

Java platform supports concurrent programming by using threads (Gosling et al., 2014). Multiple threads can execute at the same time, taking advantage of computing units that have more than one processor and of processors that have more than one core. Threads performing operations that are not atomic will interleave when they access shared data. A *synchronized* statement acquires a mutual-exclusion lock when entering a critical section, executes the block that references the shared data, then releases the lock. While one thread owns the lock for that data, no other thread may access it. The proper use of this mechanism is crucial for the concurrent implementation of an algorithm. If threads are waiting too much on locked resources, the overall performance of the application will suffer.

In our case study, multiple threads are performing operations on the file system, reading and writing large amounts of strings from and into text files. Regardless of the operating system and programming language, the underlying hardware is optimized to work with streams of bytes. In an atomic operation, data is read into a buffer of bytes, in a contiguous manner. Similarly, data is written into a buffer of bytes that is flushed afterwards to the file. Both these operations are single threaded by design and historical reasons, so Java libraries have locks to make the access to streams single-threaded. A `write` method usually looks like this:

```
public void write(String s, int off, int len){
    synchronized (lock) { ... }
}
```

We will show that creating a large number of short lived strings and writing them to a file in a multithreaded manner will generate a behavior similar to using a single thread, plus the overhead of acquiring and releasing the synchronization lock.

### 2.3 MEMORY USAGE

The Java Virtual Machine allocates memory either on stack or on heap (Lindholm et al., 2014).

The *heap* is the place where all class instances and arrays are allocated and it is shared among all threads. Each JVM thread has a private *stack* which holds local variables and partial results during successive method invocations and returns. When working with large amounts of objects, it is quite important to assess the memory consumption of a data structure, in similar way as the *sizeof* construct in C or C++. An object allocated on the heap has a *header* which contains information used for locking, garbage collection or the identity of that object. The size of

the header depends on the operating system, and it may be 8 bytes on 32 bit architectures or 16 bytes on 64 bit architectures. Also, for performance reasons and in order to conform with most of the hardware architectures, JVM will *align* data. That means that if we have an object that wraps just one byte, it will not use:  $8(\text{object header}) + 1(\text{content}) = 9$  bytes of memory on the heap, but it will use 16 bytes as it needs to be aligned to the next 8 byte boundary.

In Java, strings are objects and they are allocated on the heap. That fact that string literals are stored in a shared object pool, in order to reduce memory consumption, is of no relevance in our context. Inspecting the source code of the `String` class, one can observe the following instance fields:

```
private final byte[] value;
private final byte coder;
private int hash;
// Defaults to 0
```

As expected, a `String` object keeps a reference to an internal byte array. However, the other two fields will make the size of the object equal to:  $8(\text{header}) + 4 \text{ value reference} + 1 \text{ coder value} + 4 \text{ hash value} = 17$  bytes. Being aligned to 8 bytes, it will actually use 24 bytes. When creating many `String` instances (like millions of them, as in our case study), the extra information included in this class will add up, consuming memory and triggering the garbage collector more often than necessary.

We will show that replacing the `String` usage to the underlying `value` byte array will improve the performance of the application, and this approach should be implemented in every scenario that involves processing large amounts of text data.

## 2.4 MEMORY COMPACTION

Another important part of working with large data sets that have to be accessible in memory regards the format in which they are represented. Choosing the right format will not only reduce the amount of consumed memory but it will also reduce the GC cost to copy the objects between regions and the cost of visiting and marking them.

The most common approach of representing information is in row based form, where a *row* is a record of some kind and a *column* is a certain property of that row. This type of representation is used in most relational databases management systems, where sets of rows of the same type form *tables*. Despite having many advantages, this format is not necessarily optimum when it comes to data representation in memory.

A *column store* model (Abadi et al., 2013) "reverses" the orientation of the tables. It stores data

by columns and uses row identifiers in order to access a specific cell of the table. By storing each column separately, query performance is increased in certain contexts as they are able to read only the required attributes, rather than having to read entire rows from disk and discard unneeded attributes once they are in memory. Another benefit is that column stores are very efficient at data compression, since representing information of the same type inside of a column helps the *data alignment* process that we have previously mentioned.

Let's consider a simple example, using the class `Point`, defined as a pair of two integer fields `x` and `y`. The basic idea is that instead of having a row-based model consisting of an array `Point []` of instances (each `Point` object is a row and its members `x` and `y` are the columns), to use two arrays of integers `x []` and `y []`, representing the two columns. This way we can store the same data, minus the object headers corresponding to all the `Point` instances. Not only the memory consumption will be lower (so the GC will be triggered less often), but the structure will also take shorter time to visit, since there are only two objects now (the two arrays).

Though a column store is a very good solution for size reduction, it has the downside of requiring more computational effort in order to work with multiple properties of the same object. However, when saving memory is the major concern, and especially when it comes to hundreds of GB per instance, the execution slowdown becomes far less important if we can achieve significant reductions in consumed memory.

## 3 REPRESENTING THE DATA STORE

### 3.1 THE ROW-BASED MODEL

The data structure which is used in original EIPrep algorithm is represented by the class `SamAlignment`, an object of this type storing one row of a SAM file. The class contains the following declaration of instance variables:

```
public Slice QNAME;
public char FLAG;
public Slice RNAME;
public int POS;
public byte MAPQ;
public Slice CIGAR;
public Slice RNEXT;
public int PNEXT;
public int TLEN;
public Slice SEQ;
```

```
public Slice QUAL;
public List<Field> TAGS = new ArrayList<>(16);
public List<Field> temps = new ArrayList<>(4);
```

For a small BAM file of 144 MB there will be created around 2.1 million `SamAlignment` instances and for a 1.27 GB BAM file there will be created around 17.6 million objects.

For simplicity, let's disregard `TAGS` and `temps` fields (which can have different lengths) as it makes the calculation simpler and analyze the memory consumption in both cases. We suppose also that the JVM uses 32 bits for representing an object header.

One `SamAlignment` object contains: 8 bytes object header, 6 instances of `Slice` objects (`QNAME`, `RNAME`, `CIGAR`, `RNEXT`, `SEQ`, `QUAL`) of 4 bytes each, 3 integer fields (`POS`, `PNEXT`, `TLEN`) of 4 bytes each, 1 character (`FLAG`) of 2 bytes, and an additional byte (`MAPQ`). So, the total size of the object is:  $8 + 6 * 4 + 3 * 4 + 1 * 2 + 1 * 1 = 47$  bytes, and as it is rounded up to a multiple of 8, the result is 48 bytes.

In order to save memory, the string representing a row scanned from the original file is shared between multiple objects. All 6 `Slice` instances contain a reference to the underlying string and two integers pointing to the start index and length. So, a `Slice` instance uses: 8 (object header) + 4 (reference to the string) + 4 + 4 = 20 bytes, being rounded to 24.

As `Slice` instances point to a `String` object, the `String` itself adds another 24 bytes, as we have already seen, and the byte array object referenced from the `String` adds another 24 bytes (not counting its content size).

Adding all these numbers up, we conclude that for representing a `SamAlignment` object, the JVM needs: 48 (the object itself) + 6 \* 24 (`Slice`) + 24 + 24 = 240 bytes.

For a 144 MB file there are 2.1 million entries, so the memory requirement for storing the graph of objects and the integer fields is approximately 504,000,000 bytes, which equals to more than 480 MB (not counting the 144 MB of content in byte array).

For the 1.27 GB BAM file, the numbers are much larger as there are a 17.6 million rows. The total is 4,224,000,000 bytes, representing almost 4 GB.

At the point when GC executes, there are 9 objects per row (`SamAlignment` + 6 `Slice` + 1 `String` + 1 byte array) on heap.

The `ElPrep` algorithm will duplicate all the `SamAlignment` s once, as it looks for duplicate objects, so the required memory will double. On top of it, there are the underlying strings which keep the data. For 2.1 million entries the grand total is around 700,000,000 bytes, and 6,050,000,000 bytes for the

1 GB entry.

## 3.2 THE COLUMN-BASED MODEL

Let us analyze how much memory can be saved by switching to a column-based approach. We have defined the following data structures: `StringSequence` for representing in a compact manner a collection of strings, `DeduplicatedDictionary` for eliminating duplicate copies of repeating strings, `DnaEncodingSequence` for storing A,C,G,T,N sequences using an encoding of 21 letters per long and `TagSequence` for representing tags encoded in an array of short values. We have also used the classes `CharArrayList`, `IntArrayList` and `ByteArrayList` from `FastUtil` library (Vigna, 2019), which offers implementations with a small memory footprint and fast access and insertion.

The new definition of the data store is described by the class `SamBatch`, containing the following members:

```
StringSequence QNAME;
CharArrayList FLAG;
DeduplicatedDictionary RNAME;
IntArrayList POS;
ByteArrayList MAPQ;
DeduplicatedDictionary CIGAR;
DeduplicatedDictionary RNEXT;
IntArrayList PNEXT;
IntArrayList TLEN;

DnaEncodingSequence SeqPacked;
StringSequence QUAL;
```

So, instead of having a large number of `SamAlignment` instances, we will have a single object of type `SamBatch` which contains references to the "columns", i.e. our data structures holding all the information of a specific type.

Regardless of VM bitness, the memory consumption for representing one row of the input file is:

Data Type	Count	Bytes
<code>StringSequence</code>	2	4
<code>DeduplicatedDictionary</code>	3	4
<code>IntArraySequence</code>	3	4
<code>CharArraySequence</code>	1	2
<code>ByteArrayList</code>	1	1
<code>DnaEncodingSequence</code>	1	4
<b>Total</b>	<b>39</b>	

Considering that no rounding up is necessary, for 2.1 million rows this sums up to 81,900,000 bytes, equivalent to 78 MB. The header sizes of the column objects (11 \* 8 bytes) become negligible in this context.

*of just 60 MB of indices and after that we have the joined data.*

As this is a major reduction in memory consumption, let us analyze the technique used to achieve this result.

The basic idea is that instead of storing an array of `String` objects, for example:

```
String items[] = { "abc", "def" };
```

each consuming memory due to their headers, we can use a single object of type `String`, storing all the characters, and an additional array for their lengths.

```
String dataPool = "abcdef";  
int[] endLengths = {3, 6};
```

For such a small array, the save is minor, but for a large number of items (millions), the memory reduction becomes significant. Even more important, the GC work is also reduced, since no matter how many items are in the `dataPool` and `endLengths` fields, there are only two objects to visit. The technique described above was implemented in the class `StringSequence`.

If the strings that are to be stored are repeated frequently, we can apply another optimization: instead of keeping them joined, we will use an indexed collection containing all the distinct strings and an array holding one index for each string. For example, { "abc", "def", "abc", "xyz", "abc" } becomes:

```
table : {abc=>0, def=>1, xyz=>2}  
items : [0, 1, 0, 2, 0]
```

The `table` structure is based on the class `Object2IntOpenHashMap<String>` from `FastUtil` library, instead of the standard `HashMap<Integer, String>`, since it uses the primitive data type `int` for representing the keys, which also saves some memory. This data deduplication technique (He et al., 2010), (Manogar and Abirami, 2014) was implemented in the class `DeduplicatedDictionary`.

When storing strings containing characters from a restricted alphabet, one optimization that can be performed is using an array of primitive values, for example a `long []`, and encoding each character into a block of bits. The number of bits required for a character depends on the size of the alphabet. DNA sequences use four letters A, C, G, T, but it is possible for a machine to read incorrectly a symbol and to return N. In order to represent 5 possible characters we need at least 3 bits, which means that a `long` can store in its 64 bits 21 DNA letters. The class `DnaEncodingSequence` which implements this string encoding technique contains the following members:

```
LongArrayList content;  
ShortArrayList lengths;  
IntArrayList positions;
```

For example, encoding the 21 letters string "AAAACCCCGGGGTTTTNNNNA" would produce a single long value, containing the bits:  
00001001001001000110110110110100  
10010010001001001001000000000000

From right to left, 000 represents A, 001 represents C, and so on.

In the sample files, one DNA sequence is typically around 100 letters, so the memory needed in order to represent it would be 1 `int` (encoding length) and 5 `long`s (the content), that is 44 bytes. This reduces the memory consumption by a factor of two.

Another advantage of using such an encoding is that when checking if two sequences are exactly the same, we can compare first their lengths and, if they are equal, comparing `long` values means comparing 21 characters at once. As before, the GC will also benefit from the reduced number of objects that must be visited.

Running the smaller input file (144 MB), we have estimated that the original `EIPrep` algorithm would use around three times more memory than the size of the input SAM file. However, when trying to process the larger input file (1.2 GB), on a 32 GB machine, we have obtained an `OutOfMemoryError`, meaning that the penalty of using too many objects in order to represent the information was preventing us in loading the entire data set into memory.

#### JVM using how much memory?

1 GB input: Original algorithm: 32 024 799 640 Bytes measured in VisualVM after GC was called and program suspended right before writing to disk. This is the lowest maximum memory consumption for 1.2 GB BAM file  
Optimized algo: 31 463 351 904  
Compact algo: 4 703 674 464 bytes  
Live data memory usage: 6.8x less.  
144 MB input: Original algorithm: 2 234 575 120 Bytes measured in VisualVM after GC was called and program suspended right before writing to disk. This is the lowest maximum memory consumption for 1.2 GB BAM file  
Optimized algo: 2 273 444 096  
Compact algo: 607 589 728 bytes  
Live data memory usage: 3.67x less.  
8 GB input: 20 334 009 424 B. Obviously we had a too small machine to process a live data set that is using as minimum 3 12 GB input: 35 401 115 712 B.

Time is 300 seconds (on 8core machine) as the machine touches Swap file, but being able to process 12 GB BAM file on a 48 GB machine is a feature in itself.

With memory compaction, the memory needed to represent all the objects was estimated to 17.6 million

rows \* 39 bytes per row = 686,400,000 bytes (655 MB), plus the DNA sequences 17.6 millions \* 44 = 774MB.

As the string data was estimated to at most 3 GB, in 4.4 GB we have stored both the entire SAM file and the fields associated with the alignments.

Duplicating all this (as elPrep does it) will make the entire live set go to 9 GB, which is still around half of the original live set (what is the size of the original live set? definition of "the live set").

On top of this, there is the cost of tags. In the original implementation, every `SamAlignment` object has references to `temps` and `TAGS` arrays. These arrays have an initial size of 16, respectively 4, and contain references to `Field` objects, which in turn contain a reference to a `Slice` object. Using the row-based model, for every row the memory consumption is: 20 references \* 4 bytes + 2 \* 64 (an empty `ArrayList` costs 64 bytes) = 256 + 80 bytes. The memory usage for tags/temps is 5.9 GB. Tags only will be 80% of usage = 4.2 GB

In order to address `temps` and `TAGS` we have implemented the `TagSequence` class, which is a combination between `StringSequence` and `DeduplicatedDictionary`. Since the tags are repeating frequently, we save them using one short value per tag, but instead of using a list of tags, we define a sequence of indices.

`TagSequence` does memory compaction by using a mix of short per-tag encoding and a full sequence of tags is joined together.

For example, for the input "tag0 tag1 tag2", "tag1 tag2 tag3", the representation would be:

```
table = {"tag0":0, "tag1":1,
         "tag2":2, "tag3":3},
lengths: [0,3],
tagSequence = [0,1,2,1,2,3];
```

For a 20K rows, we found following numbers: Unique tag values: approx 2400 Total number of tags: approx 200000 (10 tags per row on average)

Numbers were correct, it was simply looking into some random batches when debugging.

## 4 OPTIMIZING I/O OPERATIONS

The first and most simple technique that could improve performance when working with large files is I/O buffering (Oaks, 2014). Invoking a Java `read` method, for example, will eventually trigger an invocation of the operating system method responsible with reading data from the disk. Although most operating systems use I/O buffering

by default, an invocation of a method outside JVM is nevertheless expensive. Another reason to minimize the number of I/O operations is that, for each one of them, temporary objects needed to store the data must be created, meaning that GC will have more work to do when recovering the memory.

The original code that we are analyzing uses correctly the `BufferedReader`, respectively the `BufferedWriter` classes in order to perform I/O operations on the SAM file. Checking the definition of the `BufferedReader` class we notice that the default buffer size is set to 8192 characters, which is reasonable for most scenarios. The `read` operation is implemented in a concurrent fashion, using Java parallel streams:

```
var alnStream = inputStream.parallel()
    .map((s) -> new SamAlignment(s));
```

In this context, for larger files (as in our case), and in the presence of multiple cores, increasing the size of the character buffer may offer a better performance, estimated by our experiments at around 5%. Este corect? Pe laptop-ul meu nu obtin nicio imbunatatire

The same simple optimization can be made for writing, by presizing the `BufferedWriter` object to a more suitable value than the default one.

When it comes to writing, we notice another possible improvements. The following code sequence iterates through the `SamAlignment` objects found in the input stream and, for each one of them, it creates a string representation containing the fields of the `SamAlignment` class and tabs as separators. This is implemented using a `StringWriter`, which is a character stream collecting its output in a memory buffer.

```
var outputStream =
    alnStream.parallel().map((aln) -> {
        var sw = new StringWriter();
        try (var swout = new PrintWriter(sw)) {
            aln.format(swout);
        }
        return sw.toString();
    });
```

The initial size of the `StringWriter` internal buffer is by default 16 characters. A line in the output file is more than 200 characters. By presizing the buffer to a value of 256, we have obtained a speed-up of the writing process estimated by our experiments to more than 10%.

If we take a closer look at the `StringWriter` class we notice that it uses an internal `StringBuffer` object. To quote from its documentation, a `StringBuffer` is a thread-safe, mutable sequence of characters. That means that most of its



methods are declared as *synchronized* in order to control the access to the buffer of characters in a multithreaded environment. However, in our case, each writing thread spawned by the parallel stream implementation gets his own copy of a `StringWriter` so there is no resource contention that would require synchronization. So, instead of using a `StringWriter`, it would be better to simply use a common buffer, with no synchronization. To maximize the performance gain, the buffer could be implemented as a byte array. Using bytes instead of characters is quite trivial as Java `String` class has the method `getBytes()` and recreating a `String` object can be done by simply using one of its constructors: `new String(bytes)`.

Writing to the file is implemented by simply invoking the `println` method of a `PrintWriter` output stream, decorating a `BufferedWriter`, which in turns decorates a `FileWriter`.


```
var out = new PrintWriter(
    new BufferedWriter(
        new OutputStreamWriter(output)));
...
outputStream.forEachOrdered(
    (s) -> out.println(s));
```

Again, the `PrintWriter.println` method is using a `synchronized` block in order to fulfill its task, which is eventually an invocation to `BufferedWriter.write` method.

```
public void println(String x) {
    synchronized (lock) {
        print(x); println();
    }
}
```

The `BufferedWriter.write` method is also `synchronized` and so is `FileWriter.write`, which is the last one invoked in this chain.

Since more than half of the I/O time is spent writing, we have analyzed the CPU utilization, using the operating system profiler.




The graph looks like a chainsaw, which is a sign that the code is bottle-necked on a single thread. This is no surprise, since the `outputStream` is written in a sequential fashion. On top of that, for every row there

has to be a mutex check whenever a `synchronized` block of code is invoked. Since there is a large number of rows, all these checks will add up and produce a significant slowdown of the overall writing process.


The technique that should have been used in this context involves using a single byte array. Instead of creating and storing individual string representations of all the `SamAlignment` objects, and writing them one by one to disk, we add progressively information into the buffer. Only when the buffer becomes full, we write its content on disk. Obviously, all these steps will be performed in a single thread, but the bottleneck will be less obvious. The class implementing the byte array is called `StreamByteWriter`.

```
StreamByteWriter streamByteWriter
    = new StreamByteWriter(output);
for (SamAlignment aln: alignments) {
    aln.formatBuffer(streamByteWriter);
}
```

Even if we are using now a single core for all the writing, the necessary time for creating the output file was reduced by half. The following graph shows the CPU utilization for this approach:



Here is the execution profile when the Original algorithm has plenty of memory and the "Chainsaw" in execution is in fact that we ran for 50 times the algorithm, and every time the CPU jumps from original 12 core usage when reading to back to 1 Core used and back. This happens only when executing on "big machines" (when the RAM available is much more than the live set, like 3-4 times at least the size)



## 4.1 CHUNKING-BATCHING AND EXTRACTING PARALLELISM

In the `elPrep` original algorithm, a `SamAlignment` object is created for each row in the input file. In order

to extract parallelism, there is an explicit `.parallel()` call, a construct that would create a task for each row that must be processed. These tasks are queued and executed in a concurrent fashion using threads created transparently by the Java Streams Framework.

```
var alnStream = inputStream.parallel()
    .map((s) -> new SamAlignment(s));
```

In the column-based model, described in section 3.2, there is a single `SamBatch` instance storing all the data. Because we perform data compaction, the *read* operation must take into account the dependencies between rows. Just like in the case of removing duplicates, in order to process a new row we have to inspect the values of the previously read rows. This means that the reading process cannot be fully parallel per-row. In order to obtain a better performance than reading using a single thread, we propose a technique that splits the `SamBatch` structure in several "chunks". Instead of having a single large object, we will represent the data using an array of smaller `SamBatch` objects. A sketch of our *Compact* algorithm is described by the following pseudo code:

```
int chunkSize = 20000;
int nbOfThreads = cpuCoreCount * 4;
var samBatches = new ArrayList<SamBatch>();
while (!endOfFile) {
    var rowsRead = parallelReadRows(
        batchSize, nbOfThreads);
    parallelTransformRowsIntoBatches(
        samBatches, rowsRead);
}
```

The `readRows` method reads `chunkSize * nbOfThreads` rows out of the SAM file for the next processing step. The actual reading is done using an appropriate number of threads (based on the CPU configuration), each thread reading sequentially a fixed number of rows, calculated taking into account the nature of the information being encoded. Having the data split into chunks, we can process in parallel the mapping between the associated text and the `SamBatch` data structure, where we perform data compaction and deduplication. The technique of grouping similar tasks requiring the same resources in order to streamline their completion is called *batching*.

In our case, the advantages of the chunking-batching approach are multiple:

- Transforming data, deduplicating and DNA encoding are all executed in a single-threaded manner, which is much easier to debug and understand than a multithreaded equivalent.
- Since a `DeduplicatedDictionary` will now have less than 20000 unique strings, the values

needed to encode the strings could be represented on 2 bytes, instead of 4. Similarly, we can reduce the tag size from 4 to 2 bytes.

- After the algorithm is fully executed, we will have instead of 2.1 million `SamAlignment` objects, about 105 `SamBatch` instances, each of them having around 4 orders of magnitude less objects overall. This translates into 2 orders of magnitude less objects.
- When creating the output file, the `SamBatch` array could be processed in parallel, with no blocking except the actual operation of writing to disk.

We have seen that the column-based model saves memory at the expense of the running time. This optimization, however, reduces the overall execution time of the read operation, which is now on par with the original implementation.

When it comes to writing, compared to our `StreamByteWriter` class, which is anyway much faster than the original implementation, the execution time is drastically reduced from 28 seconds (for the 12 GB BAM file) to 12 seconds. Preparing the strings that are to be written in the file can be done in almost perfect parallelism, using the available cores. The only limitation remains the speed of the output device, which can vary depending on its type: "in memory virtual partition", SSD, HD, etc.

In order to make sure there are no dead times when using the external device, we have also implemented the *async/await* pattern. This allows the program to perform in advance reading operations, using a dedicated thread, while waiting for the data processing threads to complete their execution. This new algorithm, called *Compact/Par* offers a small improvement in the running time, as we will see in the next section. However, the increase in code complexity makes this optimization hardly worthy of taking it into consideration.

## 5 EXPERIMENTAL RESULTS

The computer we have used in order to perform the experiments is a Ryzen 9-3900X, having 12 cores and using 48 GB of RAM. Since we didn't have access to the hardware necessary to run all the tests in memory, as the original paper, we have used the smaller SAM files and ran the same processing repeatedly in order to obtain an accurate result of the running time. For example, running 10 times the algorithm on the smallest input SAM file, which is approximately 700 MB, will produce a total running time of around 70 seconds. Running the algorithm



repeatedly will trigger the garbage collector and this will be the cause of variations in the collected results, ranging from around 2 to 3 seconds, when reading the smallest file, and 3 to 4 seconds for writing.

Before executing the timed invocations, we have *warmed-up* the JVM. This is usually achieved by simply running a couple of times an initial test (not timed) that uses all the classes involved in the algorithm. Such a warm-up is necessary because Java is using a lazy class loading mechanism and just-in-time compilation. After this step, all important classes are stored into the JVM cache (native code), making them available at runtime with no additional penalty.

The original elPrep benchmarks have been performed on a Supermicro SuperServer 1029U-TR4T node with two Intel Xeon Gold 6126 processors consisting of 12 processor cores each, clocked at 2.6 GHz, with 384 GB RAM (Costanza et al., 2019). The authors claim to do the processing of the 8 GB BAM file in 6 min:54sec to 7 min:31sec and memory usage is 330 – 340 GB.

As we didn't have access to such a performing machine, we did most of testing with the smallest file, the 144 MB BAM file (673.3 MB SAM file). For the 8 GB BAM file (27.18 GB SAM) our results will show only the compact algorithm but we will make some inference over the scaling of the original algorithm across file sizes and scalling over cores.

Especially as elPrep algorithm was using a lot of memory compared to Compact algorithms, we tuned the JVM heap size (using `-Xmx` flag) to maximum value that the OS would run everything in memory, which makes in fact numbers to be a bit better than using just the defaults. Without it, values like 123 seconds for original algorithm would degrade into either around 200 seconds mark for runtime or Out-of-memory errors.

For all algorithms we did use the fastest disk possible (like we were choosing an SSD to read, and if possible, we would be using either RAM or fastest spinning disk). And as disk access is sometimes affected by antivirus service in Windows (and as we tested on a regular Windows 10 machine) or background services like Dropbox, OneDrive, etc all services that could be noticed would be manually stopped, and made sure that there is no other service running and affecting the numbers (like Windows Updates). To try to make sure of no extra factors were affecting performance, Windows Task Manager was running in the background, but it was not found to affect the performance of the tests, especially on the 12c machine.

## 5.1 RUNTIME PERFORMANCE

### Algorithm times

	0.15 (4c)	0.15 (12c)	1.2 (12c)
Original	9.13 s	3.938 s	123.91 s
Presized buf	8.98 s	4.19 s	75.1 s
<i>StreamByteArray</i>	8.09 s	3.43 s	64.62 s
Compact	3.16 s	1.686 s	34.5 s
Compact/Par			26.7 s

Analysis of speedup across machines (4 core vs 12 core): Original algorithm scales with a factor of 2.3, presized buffers by a factor of 2.14, *StreamByteArray* scales by 2.36 and compact would scale by: 1.686, so at least for small files it looks that scaling over large machines will be able to offer a better scaling. Also, it is important to notice that the compacting algorithm has an explicit sequential part that reduces scalability, and it may be reduced with proposed ideas of potential improvements.

Scaling across file sizes: the original algorithm uses a lot of memory, and as the VM has to call more often the Garbage Collection, the cost of GC affects visibly the scalability of the original algorithms. As GC happens in the Compacting algorithms, the cost of it doesn't affect runtime and reducing allocations in both of the techniques of "Presized Buf" and "StreamByteArray" do show a much better runtime profile:

So for an increase of input of 8 times, the original algorithm would slow down 31.46 times, the "Presized buffers" will slow down 17.92 times, "StreamByteArray" will slow down 18.83 times, but Compact algorithm will have almost a perfect 8.33 times slow down.

As times will grow for a 8 GB BAM file to around 215 seconds (when setting heap size to be 42G as maximum using `-Xmx42G`), we can investigate and isolate the sequential I/O cost. So as we are reaching the limits of the machine, for a file that is approximately 5 times larger, the algorithm slows down by a factor of 12. When moving the input file on an SSD, the times improve to 139.8 seconds for compact algorithm. If moving to a bigger batch size (from 20K to 30K), the algorithm will process the output in 133.15 seconds, but for the sake of comparison, it is possible to conclude that using fast I/O can improve the runtime performance so the slow-down is from 12 times to 10 times (from 14.08sec to 140 sec).

## 5.2 MEMORY USAGE

To measure the live data set, we use the already existing instrumentation in the JVM and the

visual tool VisualVM to capture according to different algorithms the different values of memory consumption and it describes the minimum amount of memory that JVM has to have available theoretically to load specific data set. To get the values from the following paragraphs, it was clicked on the tab "Monitor", look into "Memory" section and we triggered forcefully a GC after we suspended the algorithm when loading all the model (which is the maximum memory used) and it was all used for reading.

Compared with elPrep paper which measure process memory size: we measure live data size, not per process memory usage.

Because of VisualVM, we can get a very precise memory consumption and many variables are excluded like: the G1GC (the default GC algorithm in Oracle JDK/OpenJDK) uses background threads to free memory which interact with threads of execution, having a quite random values on runtime.

It is also important to understand that the more of live memory is used and the more complex the graph of objects is done and the less free memory exists on a particular machine, the GC impact will be higher. A hypothetical GC algorithm could be triggered more seldom if it has double the free memory, because the full heap space necessary to trigger GC is occupied more seldom.

As for defaults: G1GC would automatically trigger according to documentation at 45 percent memory usage mark, and when the heap is arriving at this level of occupancy, will span background threads to do the marking phase, so having a larger machine with double of RAM would trigger GC at least 2 times as seldom, even the GC time would take similar times given the similar number of objects (According to docs: -XX:InitiatingHeapOccupancyPercent=45)

Live data set would let us avoid some of the noise given that most GC algorithms do have typically a "saw teeth" profile, making hard to say how much memory is really used.

**Memory usage per algorithm in MB**

Algos	144 MB BAM	1.2 GB
Original	2 326 MB	32 025 MB
<i>StreamByteArray</i>	2 275 MB	31 463 MB
Compact/Par	606 MB	4 689 MB

Compact algorithms can process on a 48 GB machine larger files like the 8 GB BAM file and uses: 24 283 MB of live data and the 12 GB BAM input is using 35 265 MB of memory.

Even the live memory set is showing the peak memory usage (by suspending the program and measuring when the whole file is read), on practical terms all algorithms just before finishing (let's say at 99.9 percent mark of finishing reading), would use a

bit more live memory but not too much and it is quite hard (if not impossible) to measure it. For example in case of elPrep, as it creates millions of tasks even for smallest 144 MB BAM file, Java Streams library will create a queue of tasks that will not be freed very likely until the very end, adding probably few tenths of maybe 100 MB of memory. Similarly the compact algorithms would keep in memory last set of chunks (a batch) that would be similarly even on a large machine under 100 MB and they will be unused right after the related strings are converted into data.

For all practical purposes, the numbers here are very precise and they would show the range of the machine that could process that data set.

So for example elPrep "Original" or "Presized Buffers", would use 32 GB of memory at peak usage (in case of 1.2 GB BAM file), so to process it efficiently, probably a machine with 48 GB (with heap tuning, as we did to get best numbers for elPrep, by using -Xmx flag in JVM) or a machine with 64 GB would suffice.

## 5.3 CALCULATING PERFORMANCE X MEMORY


There are limits on the testing that could be done with the techniques used by this paper given the hardware differences that are between the elPrep paper and this one: instead of a 384 GB of memory and dual-CPU server, the most powerful machine that was available was a 8 times less RAM, with half of CPU cores, a much slower disk drive (the SSD on the Ryzen machine can barely read at a 300 MB/s, vs elPrep paper which had very likely an NVMe drive, that is much faster) and so on.

Another limitation is that we ported only the most expensive part of the algorithm (reading/writing and storing all in memory) and we will have a section with possible untested improvements that could be done to get (at least it looks reasonable that) better numbers that are obtained compared with the results that are presented here.


So given that memory remains quite similar between: original algorithm, pre-buffered one and using a StreamByteArray custom class, the runtime savings should be observed distinctively by noticing the runtime savings.

On very small files, there is a large variation between runtimes as they run just in seconds.


The high core count (12 c) has a small slow-down running for 100 times the parallel algorithm which is reflected on the graph (lower bars are better):



On a laptop, because of very different components (including a better SSD), the compact algorithm runs faster (lower bars are better):




On the medium file (1.2 GB BAM file) the times are very decisive, and they should reflect (given the difference in scale of the machines are tested between a 48 GB RAM and the original paper that was using 384 GB of RAM) the result scaling on this paper (lower bars are better):



It is important to see if given a larger file, if the GC becomes a large impediment and up-to the 8 GB file, the algorithm seems to scale nicely and to be limited by other components, like the SATA controller speed and the disk drive itself (or maybe OS components).

Both compact algorithms (Compact and Compact-Parallel) keep a steady pace or read/write of around

180 MB/second (averaged and it includes both read/write, higher bars are better):



If the running machine would include a very fast hard-disk, probably through profiling some other component could become a slow-down (like the rate of allocation in the JVM), but very likely the ratio of described Compact algorithms would remain quite steady up-to the time when the JVM/OS will remain without RAM, but as the algorithms use a typically a 4-6 times memory savings compared with elPrep algorithm, the limits would be reached much later. If limits like the ratio of allocation of JVM becomes the limit, few proposed techniques are described in section "Future work" that may address this limitation.

## 5.4 OVERVIEW AND ANALYSIS OF RESULTS

We found 5 underlying algorithms that all do some levels of improvements in runtime performance, especially on large files as GC becomes the limiting factor:

- \* the original algorithm described in original implementation (123.91 sec)

- \* "Presized buf" which avoids quite many allocations by presizing on writing the expected size of the row, eliminating many allocations and especially if the machine's heap reaches the limits of the machine can give a huge speedup (75 seconds)

- \* "optimized" algorithm which doesn't technically speed up the algorithm, but it makes the writing to do almost zero allocations and to write directly the bytes into a OutputStream. It was achieved without reducing debuggability (meaning that if suspending execution to look for a bug, toString method displays the same relevant data)

- \* isolating the slow parts of the original algorithm as being memory heavy and IO heavy (Read-write) we can achieve a comparable performance on read/write, even there is a computational cost on every

element, with an observed 4 times less live-memory (34.5 seconds)

\* we can isolate even more a compact version of algorithm (named "Compact parallel") that as it compacts data, it also does prefetching for IO (it reads in a background task the next block that will be needed to be compacted on the next operation), and the wins are visible on slow disk, but even using a RAM disk, the algorithm goes even faster (33.38 sec per iteration).

## 6 FUTURE WORK

### 6.1 SCALABILITY OF BATCH READER

Our reader works as described before in two steps:

\* it separates initially from the original file the rows for the number of the expected batches on the main thread

\* it executes in parallel the compaction step over the separated rows over the all cores

Speeding up the single threaded code by not making allocations

For now, before chunking and batching is done, we split the full byte array read from file into discreet `List<byte[]>` instances, but this construct may not be necessary if it is found to be a bottle neck: an alternate approach it would be to store inside the large `byte[]` buffer read the

`n` character positions, so no splitting is necessary, but it would necessitate that the reading inside the line to be rewritten.

This type of coding that would strongly avoid almost any allocations would be necessary just if the memory allocator is a bottleneck in runtime execution.

A similar variation of speeding up the single threaded part is to not do splitting into lines at all, but to read an approximate size in advance that would be around the expected chunk size and split it after in lines in a multi-threaded way.

### 6.2 TESTING THE CODE USING VALUE TYPES

As memory consumption and allocation seem to be large slow-downs inside `elPrep` and it is having an impact on very large data sets also on the algorithms which we name them Compact, reducing even by a bit the memory usage and memory compaction. .Net's Virtual Machine named CLR

has a very compact memory implementation of `HashMap` (`System.Generic.Dictionary` class) that could be possibly giving even more memory savings. Similarly, there is an effort named Project Valhalla in JVM that has previews to download that can technically implement a memory compact dictionary that may have the same memory savings.

So just porting to .Net the code may bring less memory usage and maybe some speedups.

## 6.3 DOING THE ALGORITHMS

When implementing the Compact algorithms, it was implemented only the read/write and materialization of the data set stored in SAM/BAM files, as it was a very expensive part of the `elPrep` implementation.

Though the data in-memory has to be reordered, filtered, deleted from invalid data, and so on. These algorithms very likely can be very easily made in parallel (as any chunk is a parallel subset of the entire file), but it isn't both the point of this paper as it focused only on memory usage and very fast reading/writing.

Some notes can be easily expected: \* some tasks can be integrated into reading before inserting data into `SamBatch` structures, but it will be at a cost of code complexity, so a more sensible approach would be to read first, then do algorithms one by one in a parallel fashion \* filtering and deleting data in general will have to include two steps (or a secondary `SamBatch` which rebuilds from the first only with items that pass the condition) \* sorting is possible as the batches are directly indexable, so item with number: 654321 would be on `SamBatch` index  $654321 / 20\,000$ , and with sub-index inside the batch of  $654321 \% 20\,000$  \* another way to do the sorting is to sort inside the batch first by sequences (and sequences are easily sortable as the sorting is done 21 DNA letters at once) inside the batch, then a merge sort could be done across the `SamBatch(es)`

## 6.4 CONCLUSIONS

The presented techniques in this paper show that in algorithms that are heavily limited by memory allocation(s), just by reducing creating many strings can improve the runtime performance. We measured a speedup factor of 2 in large files just by writing sequentially and not allocating anything as we write the bytes ourselves by tuning `ElPrep` algorithm.

As `elPrep` stores a large data-set looking like a large CSV file - and for similar algorithms that need to store large data sets in memory - techniques of memory compaction, especially tuned per type of

column, can reduce significantly the memory usage without reducing the performance of the application, compaction time being typically less than calling underlying garbage collection.

If memory compaction is mixed with chunked data and batching the chunks are taken proportional with numbers of physical cores, the algorithm can extract parallelism making the code to be comparable with implementations that extract good parallelism as is elPrep.

Using parallel fetching as the algorithm of memory compaction is happening (algorithm named by us "Compact Par"), it can improve even more the performance for these algorithm (from 159 secGB to 125 secGB).

## REFERENCES

- Abadi, D., Boncz, P., and Harizopoulos, S. (2013). *The Design and Implementation of Modern Column-Oriented Database Systems*. Now Publishers Inc., Hanover, MA, USA.
- Costanza, P., Herzeel, C., and Verachtert, W. (2019). Comparing ease of programming in C++, Go, and Java for implementing a next-generation sequencing tool. *Evolutionary Bioinformatics*, 15:1176934319869015.
- Gosling, J., Joy, B., Steele, G. L., Bracha, G., and Buckley, A. (2014). *The Java Language Specification, Java SE 8 Edition*. Addison-Wesley Professional, 1st edition.
- He, Q., Li, Z., and Zhang, X. (2010). Data deduplication techniques. volume 21, pages 430 – 433.
- Herzeel, C., Costanza, P., Decap, D., Fostier, J., and Verachtert, W. (2019). elprep 4: A multithreaded framework for sequence analysis. *PLOS ONE*, 14(2):1–16.
- Java Platform, Standard Edition (2019). Java Development Kit Version 11 API Specification. <https://docs.oracle.com/en/java/javase/11/docs/api>. Accessed: 2019-06-01.
- Lindholm, T., Yellin, F., Bracha, G., and Buckley, A. (2014). *The Java Virtual Machine Specification, Java SE 8 Edition*. Addison-Wesley Professional, 1st edition.
- Manogar, E. and Abirami, S. (2014). A study on data deduplication techniques for optimized storage. pages 161–166.
- Oaks, S. (2014). *Java Performance: The Definitive Guide*. O'Reilly Media, Inc., 1st edition.
- Oracle (2019). Java garbage collection basics - oracle. <https://www.oracle.com/webfolder/technetwork/tutorials/obe/java/gc01/index.html>. Accessed: 2019-06-01.
- Vigna, S. (2019). Fastutil 8.1.0. <http://fastutil.di.unimi.it/>. Accessed: 2019-06-01.